Building the Next-Gen Legal AI Platform Together

PRESENTED BY

Marcelo Caballero, MSc Artificial Intelligence, MBA

AUDIENCE

Dr. Paulina Grnarova, CEO & Co-Founder

Dr. Yannic Kilcher, CTO & Co-Founder

Dr. Kevin Roth, COO & Co-Founder

PURPOSE

Application for Product Manager (Lead) Role

Demonstrating Strategic and Technical Fit for Next-Gen Legal AI

ORGANIZATION

DeepJudge AG, Zurich, Switzerland

Purpose and Intent

I approach this opportunity with deep respect for what you have built. DeepJudge represents a rare achievement: **a technically sophisticated platform that has achieved product-market fit in one of the most conservative, risk-averse markets in enterprise software.**

WHAT THIS PRESENTATION WILL COVER

- My understanding (based on public data) of DeepJudge's strategic positioning as foundational infrastructure, not a generative AI application

- Deep analysis of your product architecture, technology stack, and platform evolution

- Evidence of my technical and strategic fit for the Lead Product Manager role

- A forward-looking vision for the next 12 months of platform development

MY APPROACH FOR THIS CONVERSATION

- **Substance-Driven**: A focus on strategic analysis and the technical realities of building an enterprise platform, not on marketing language

- **Evidence-Based:** A pragmatic assessment of the opportunity, grounded in your public-facing strategy and my own professional experience

- **Forward-Looking**: A direct and honest discussion about the complex "dual-user" and "dual-stakeholder" challenges that must be solved to win

This is a conversation among peers who understand the technical and strategic realities of building AI platforms for enterprise markets.

Why This Opportunity Matters

The convergence of three strategic inflection points

1 **Platform Maturity: From Tool to Ecosystem**

DeepJudge has completed the critical transition from search tool (2021-2023) to RAG assistant (2024) and is now entering the **platform phase (2025+)**. The AI Workflows product transforms DeepJudge from a tool that lawyers use into a platform that firms and partners build upon. This is the classic **"0-to-1" platform opportunity**.

2 **Market Validation: The Thomson Reuters Partnership**

The October 2025 partnership with Thomson Reuters validates DeepJudge's strategic thesis. TR spent **$650 million** to acquire Casetext and master AI on public legal data. Two years later, TR needed to partner with DeepJudge to access the "missing half", secure, permission-aware retrieval of internal firm data. This partnership is not a feature integration. It is the foundation of **TR's multi-billion-dollar legal AI strategy**.

3 **Personal Alignment: Platform PM at Scale**

My career has been defined by building AI and IoT platforms on top of foundational infrastructure. At Vorwerk, I managed a **CHF 100 million** connected product portfolio, building the "Thermomix as a Service" platform. At V-Zug, I delivered predictive maintenance and analytics for **500k+ connected appliances**. I have spent a decade bridging the gap between world-class engineering teams and enterprise customers with stringent security, compliance, and data governance requirements. **This role is the natural evolution of that trajectory**.

The Knowledge Layer Thesis: Infrastructure, Not Application

> **"The only AI platform that knows everything your firm knows"**

DeepJudge is not a Generative AI application. It is **foundational knowledge infrastructure**.

While competitors lead with generative AI applications (e.g., "draft this brief"), DeepJudge leads with infrastructure.

THE STRATEGIC THESIS

1	**2**	**3**
LLMs are Commoditized	**Law Firms' Data is Siloed**	**Retrieval is the Moat**
Any competitor can access GPT-4, Claude, or Gemini. The differentiator is not the model. It is the data the model can access.	Institutional knowledge is scattered across DMS, SharePoint, HighQ, email archives, and individual expertise. This data is high-value, private, and inaccessible to generic AI tools.	The founding team's expertise from Google Brain, Google AI Language, and Google Search positions DeepJudge to solve the hard problem: secure, permission-aware retrieval at enterprise scale.

PLATFORM STRATEGY: "INTEL INSIDE" FOR LEGAL AI

By becoming the "core AI platform" for a firm's entire private data universe, DeepJudge positions itself as the secure "operating system" for legal AI.
This is a far stickier, more defensible, and more valuable long-term position than selling a single-use-case application.

Knowledge Layer Thesis: Market Evidence

EVIDENCE VALIDATING THE INFRASTRUCTURE-FIRST APPROACH

GUNDERSON DETTMER VALIDATION

A sophisticated 400-lawyer tech firm built its own proprietary GenAI app ("ChatGD") in 2023, then **replaced it with DeepJudge** in March 2025. This validates the infrastructure-first approach: even a well-resourced, technically sophisticated firm concluded that DeepJudge's foundational retrieval layer was superior to building their own.

THOMSON REUTERS PARTNERSHIP

Even a **$650 million acquisition** (Casetext) could not solve the internal data retrieval problem without DeepJudge's "breakthrough contextual enterprise search." Thomson Reuters, after mastering AI on public legal data, still needed to partner with DeepJudge to access firms' internal knowledge repositories.

Strategic Implication: These two cases prove that DeepJudge's infrastructure-first approach is not a positioning choice.

It is a technical reality. The retrieval problem is hard enough that neither a sophisticated law firm (Gunderson Dettmer) nor a $650M acquisition (Casetext) could solve it independently. DeepJudge's moat is built on solving the problem that everyone else underestimated.

The Founders' Vision: ETH/Google Heritage

World-class pedigree in the exact domain where legal AI is won or lost

Dr. Paulina Grnarova

CEO & CO-FOUNDER

Background: Google Brain, ETH Zurich PhD in Generative AI

Recognition: Forbes 30 Under 30, Digital Shaper of Switzerland

Strategic Role: Vision and external-facing leader, managing key partnerships (Thomson Reuters)

Dr. Kevin Roth

COO & CO-FOUNDER

Background: Google Brain (Berlin), Microsoft Research (Cambridge), ETH Zurich PhD in ML

Expertise: Physicist turned computer scientist, data and algorithms specialist

Strategic Role: Operations and execution leader, oversees company operations

Dr. Yannic Kilcher

CTO & CO-FOUNDER

Background: Google AI Language, ETH Zurich PhD, training and deploying ML models at scale

Public Profile: 302,000-subscriber YouTube channel on ML research

Strategic Role: Technical and product architect, public AI thought leader

The founding team's combined **Google/ETH pedigree** forms the core of DeepJudge's technical and cultural DNA. All three co-founders share formative career experience as researchers and engineers at Google, specifically working on **Google Brain, Google AI Language, and Google Search,** the precise technical domains (retrieval, search, enterprise ML) where legal AI is won or lost. Their academic foundation at ETH Zurich, one of Europe's most prestigious science and technology universities, provides the rigorous research background that underpins DeepJudge's technical depth.

The Hofmann Lineage: Second-Generation Advantage

All three founders were PhD students in the **Data Analytics Lab at ETH Zurich**, led by Professor Thomas Hofmann.

Hofmann is not merely an academic mentor, he is a serial entrepreneur in this exact domain.

PROFESSOR THOMAS HOFMANN

• **Co-founded Recommind in 2000**, one of the foundational companies in e-discovery and AI-powered legal search

• **Successfully exited Recommind to OpenText in 2016**

• Serves as a **key investor and advisor to DeepJudge**

This lineage means DeepJudge is a **"spiritual successor"** to a major legal tech exit. The founders have been guided from "day zero" by a luminary who has already built, scaled, and exited a successful company in the last tech cycle for legal search. This provides an immense, "un-learnable" head start in domain expertise, enterprise sales strategy, and investor networks.

Strategic Implication: The founding team possesses world-class expertise in the precise technical domain (retrieval, search, enterprise ML) where legal AI is won or lost. This is not a team of LLM specialists building a chat wrapper. This is a team of search engineers building the knowledge layer.

Market Validation: Unparalleled Adoption Metrics

Adoption rates that differentiate DeepJudge from typical legal tech shelf-ware

CLIENT ADOPTION RATES

>80%

Homburger

Legal professionals incorporating into workflow. Partner David Oser describes this as "unparalleled" and calls DeepJudge a "cornerstone technology within our tech stack."

90%

Lenz & Staehelin

Adoption rate among legal professionals. Partner Fedor Poskriakov calls it a "Breakthrough Technology."

>70%

LEXR

Of the firm's members regularly use DeepJudge.

+23

Queries per user per day, including 10+ AI Workflows per user per day

Market Validation: Strategic Wins

THREE MAJOR STRATEGIC WINS

1. **Gunderson Dettmer (March 2025)**

A sophisticated, 400-lawyer tech-focused firm replaced its own proprietary GenAI app ("ChatGD") with DeepJudge. CIO Joe Green is now using the AI Workflows platform to build "ChatGD+," stating:

"The DeepJudge Workflow platform enables us to create AI applications tailored to our different departments."

2. **SKILLS.law #1 Recommendation**

DeepJudge achieved the #1 Net Promoter Score among all legal tech tools in the SKILLS.law 2024 survey, outperforming established players like LexisNexis, Westlaw, and Practical Law.

3. **Top 100 Swiss Startups: Ranked #3 (2025)**

Highest position achieved, validating elite credibility within the European Deep Tech ecosystem.

What these metrics prove: Deep, daily-use integration (not shelf-ware) • Firm-wide dependency (DeepJudge becomes the de facto knowledge layer) • Validation from the market's most sophisticated buyers

The Thomson Reuters Catalyst: Completing the 360° View

The partnership that transforms DeepJudge from vendor to strategic infrastructure

PARTNERSHIP SCOPE

The **October 22, 2025** announcement of a "legal data partnership" with Thomson Reuters is the single most important event in DeepJudge's history and the primary catalyst for the "Next-Gen" PM role.

> **STATED GOAL**
>
> Bring DeepJudge's AI-powered enterprise search capabilities to law firms and legal departments through Thomson Reuters trusted CoCounsel Legal

THE 360° VIEW

→ **Internal Firm Know-How:** Powered by DeepJudge's "breakthrough contextual enterprise search"

→ **External Exclusive Content:** Powered by TR's "trusted content" (Westlaw, Practical Law)

THE CASETEXT "GAP"

$650M

In August 2023, TR acquired Casetext (maker of CoCounsel) to master AI on public legal data. TR's grand strategy: be the "technology backbone of the legal industry."

The strategic gap: Casetext technology did not have the "breakthrough" capability to securely index and retrieve data from all of a firm's internal, siloed systems at scale while respecting complex permissions and ethical walls.

Result: Just two years after its landmark acquisition, TR had to partner with DeepJudge to access the "missing half"—the firm's internal work product locked in Document Management Systems.

Strategic Implication: DeepJudge is not a minor "bolt-on". It is the missing, foundational piece of Thomson Reuters' multi-billion-dollar AI strategy. This places the Lead Product Manager, the person who will build the workflows connecting these two systems, at the absolute center of the most important strategic alliance in legal tech today.

Competitive Positioning: Four Pillars of Differentiation

1

World-Class "Retrieval" Credibility

Founded by ex-Google Search engineers (Google Brain, Google AI Language, Google Search) to solve enterprise search, not by LLM specialists building a chat app. The founding team's pedigree is in the precise technical domain where legal AI is won or lost: retrieval, not generation.

2

True Data Unification

The only platform designed from the ground up to index all scattered knowledge (DMS, SharePoint, HighQ, email, intranets) into a single, unified knowledge layer. Competitors focus on single-source or manually uploaded documents. DeepJudge provides firm-wide, automated, permission-aware unification.

3

"Fort Knox" Security & Architecture

Flexible deployment (cloud, private cloud, on-premises) | No training on client data | Permission-aware indexing mirroring all access controls and ethical walls | LLM-agnostic architecture. This is a strategic weapon that opens doors cloud-only competitors cannot enter (Swiss banks, German industrials, national security firms).

4

Proven "Unparalleled" Adoption

80-90% adoption metrics at elite firms (Homburger, Lenz & Staehelin) are the ultimate proof that the product works and delivers daily value.
#1 Net Promoter Score (SKILLS.law 2024), de-risking the purchase for new clients. This is the antidote to legal tech's chronic "shelf-ware" problem.

COMPETITIVE LANDSCAPE SUMMARY

Harvey	Horizontal generalist, $3B-$5B war chest. Weakness: cloud-only, OpenAI-dependent, lacks deep retrieval expertise.
Luminance	Contract-focused specialist. Weakness: narrow M&A use case, not a platform.
Microsoft Copilot	Bundled giant, "at hand" for M365 users. Weakness: "good enough" but not best-in-class, lacks legal-specific retrieval depth.
TR CoCounsel	Strategic partner and potential competitor. Weakness: lacks internal data retrieval, hence the DeepJudge partnership.

Architecture Overview: Layers 1 & 2

DeepJudge's product evolution is a masterclass in platform strategy

---

L1 Knowledge Search

2021-2023

INFRASTRUCTURE AS A SERVICE (IAAS)

Product	Advanced enterprise search across all firm repositories (DMS, SharePoint, HighQ, email)
Technology	Proprietary hybrid search combining keyword/metadata matching with semantic vector search
Strategic Purpose	The "Trojan horse" entry point. Solves universal pain point of broken internal search. Creates firm-wide dependency.
Go-to-Market	"Land" strategy, enter the firm by solving a tangible, C-suite-visible problem
Evidence	Lenz & Staehelin and Homburger confirm this "advanced search engine" was the entry point

---

L2 Knowledge Assistant

2024

SOFTWARE AS A SERVICE (SAAS)

Product	RAG-powered chat interface built on top of the Knowledge Search infrastructure
Technology	Retrieval-Augmented Generation combining DeepJudge's hybrid search with LLM-based generation
Strategic Purpose	The first "app" built on DeepJudge's own infrastructure. Demonstrates that the search layer can power generative AI applications.
Go-to-Market	"Expand" strategy, upsell existing Knowledge Search customers to a higher-value tier
Evidence	Announced June 2024 alongside the $10.7M funding round led by Coatue

Architecture Overview: Layer 3 and Platform Strategy

The most critical layer for the Lead Product Manager role

L3   AI Workflows

2025

PLATFORM AS A SERVICE (PAAS)

Product
Low-code/no-code workflow builder enabling firms and partners to create custom, multi-step AI applications

Technology
Agentic AI orchestration layer that can plan and execute multi-step legal tasks, integrating multiple tools, models, and data sources

Strategic Purpose
Elevates DeepJudge from a tool (that lawyers use) into a platform (that firms and partners build on). Creates an ecosystem of "builders."

Go-to-Market
"Monetize" strategy, dual-user model serving both end-users (lawyers) and builders (legal engineers, partners like TR)

Evidence
Gunderson Dettmer CIO Joe Green is using the AI Workflows platform to build "ChatGD+," tailored applications for different departments

The Lead Product Manager role is the owner of Layer 3. This is not a "SaaS PM" role adding features to a chat app.

This is a "PaaS PM" role managing a platform and fostering its ecosystem of builders.

Core Technology Mechanics: Hybrid Retrieval & Permissions

The first two pillars of DeepJudge's technical architecture

1   Hybrid Search

DeepJudge combines **keyword/metadata matching** (traditional Boolean search) with **semantic vector search** (embedding-based similarity). This dual approach delivers both precision (finding exact clauses, case citations) and recall (finding conceptually similar documents even when terminology differs).

> **Strategic Advantage:** Outperforms pure keyword search (misses semantic matches) and pure vector search (misses exact matches). This is the technical foundation that enables the "unparalleled" adoption rates.

2   Permission-Aware Indexing

DeepJudge mirrors all access controls from source systems (DMS, SharePoint, HighQ) at the index level. Every document inherits its original permissions, ethical walls, and matter-level access restrictions. A lawyer searching DeepJudge sees only the documents they would have access to in the original repository.

> **Strategic Advantage:** This is the "Fort Knox" security architecture that enables DeepJudge to index firm-wide data without creating a security vulnerability. Competitors that require manual document upload cannot achieve this level of automated, permission-aware unification.

Core Technology Mechanics: Unification & Architecture

Continuing the technical deep dive: data unification and platform flexibility

3 Data Unification

DeepJudge provides a **single unified API** across all of a firm's scattered knowledge repositories: iManage, NetDocuments, SharePoint, HighQ, email archives, intranets. The platform automatically crawls, indexes, and keeps these sources synchronized, creating a single "knowledge graph" of the firm's institutional memory.

> **Strategic Advantage:** Eliminates the "search 5 different systems" problem that plagues law firms. This unification is the foundation for the "360° view" that Thomson Reuters needs to combine internal and external data.

4 LLM-Agnostic Architecture

DeepJudge's orchestration layer is **model-agnostic**. The platform can route queries to OpenAI (GPT-4), Anthropic (Claude), Google (Gemini), or locally-hosted open-source models. This flexibility allows firms to choose models based on cost, performance, data residency requirements, or regulatory constraints.
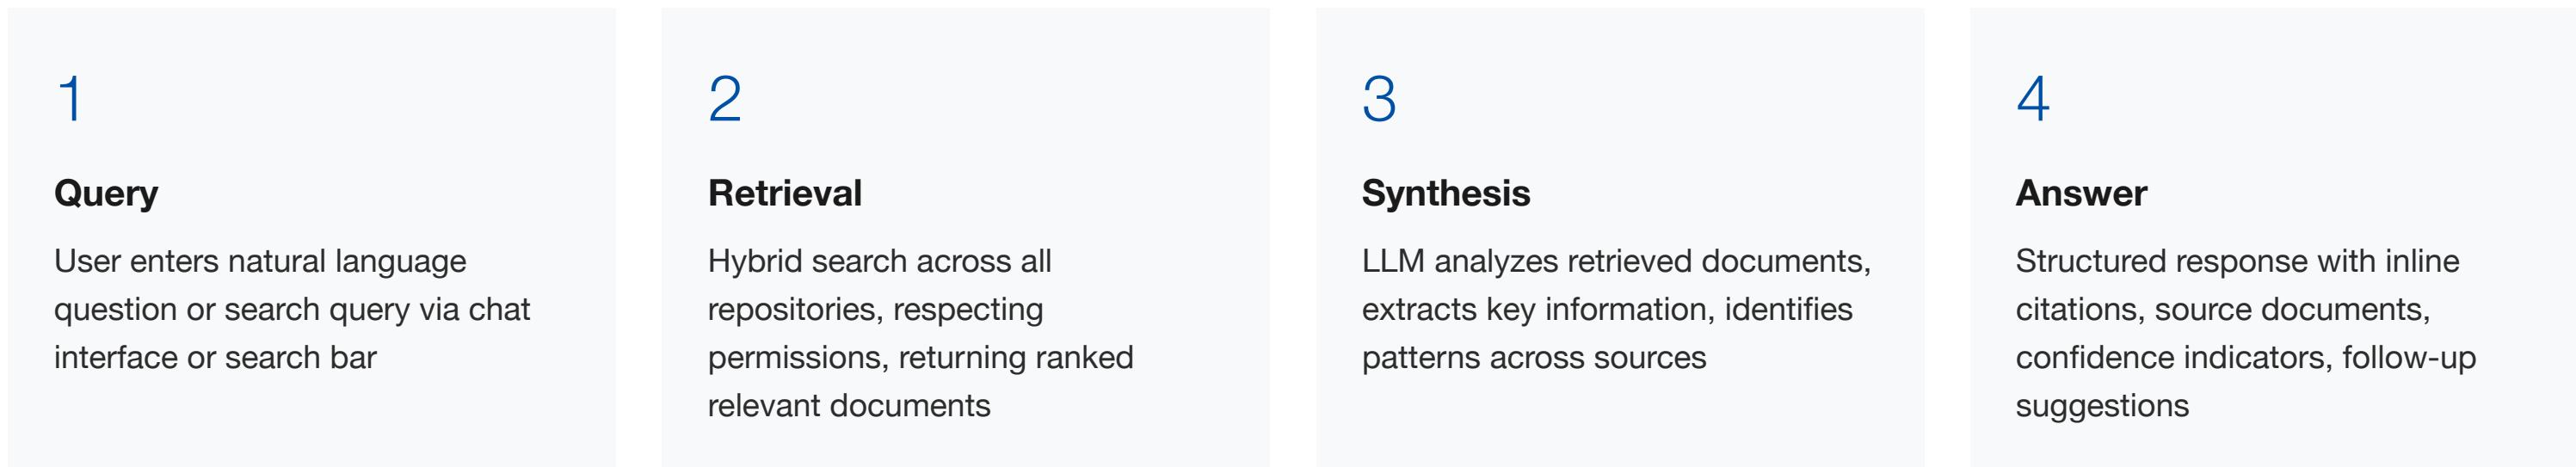
> **Strategic Advantage:** De-risks vendor lock-in and future-proofs the platform. As new models emerge (or as regulations require on-premises deployment), DeepJudge can adapt without re-architecting the entire system.

The Technical Moat: These four pillars create a defensible competitive advantage. Competitors can replicate the chat interface in weeks, but replicating permission-aware, firm-wide data unification with hybrid search requires years of domain expertise and deep integration with legacy enterprise systems.

Platform Capabilities Today: The UX Journey

END-TO-END USER EXPERIENCE

1	2	3	4
Query	**Retrieval**	**Synthesis**	**Answer**
User enters natural language question or search query via chat interface or search bar	Hybrid search across all repositories, respecting permissions, returning ranked relevant documents	LLM analyzes retrieved documents, extracts key information, identifies patterns across sources	Structured response with inline citations, source documents, confidence indicators, follow-up suggestions

MULTI-DOCUMENT INTELLIGENCE

→ **Cross-Document Analysis:** Synthesize information from 10+ documents simultaneously

→ **Pattern Recognition:** Identify recurring clauses, precedents, and legal strategies across matter history

→ **Comparative Analysis:** Compare contract terms, case outcomes, or regulatory positions

→ **Temporal Tracking:** Track evolution of legal positions or contract negotiations over time

→ **Provenance Tracking:** Every claim linked to source document with page/paragraph citations

Platform Capabilities: Current Workflow Examples

CURRENT WORKFLOW EXAMPLES

Due Diligence Checklists

Auto-generate checklists from historical M&A matters

Contract Review

Compare new contract against firm's standard templates and past negotiations

Legal Research Memos

Synthesize internal precedents with external case law

Compliance Checks

Verify regulatory compliance across multiple jurisdictions

Matter Onboarding

Surface all relevant past work on similar matters or clients

Know-How Extraction

Extract reusable templates and playbooks from closed matters

Key Differentiator: DeepJudge is not a search engine that returns documents. It is an intelligence platform that synthesizes answers from documents. The user receives a structured, cited response, not a list of PDFs to manually review. This is the difference between "here are 50 documents" and "here is the answer, supported by these 50 documents."

Next-Gen Direction: Evolution to Agentic AI

EVOLUTION TO AGENTIC WORKFLOWS

DeepJudge is evolving from **single-shot Q&A** (RAG-powered chat) to **multi-step autonomous workflows** (agentic AI). Instead of "answer this question," the platform will "complete this task" by planning and executing multi-step processes: conduct due diligence across 50 documents, compare contract terms against 10 precedents, verify regulatory compliance across 3 jurisdictions. This is the shift from **assistant to agent**.

THE DUAL-PERSONA CHALLENGE

Users (Lawyers)

END-USER PERSONA

- Want pre-built, curated workflows

- Expect simple, intuitive UX

- Need immediate, reliable value

- Low technical fluency

- Measured by adoption and daily usage

Builders (Legal Engineers / Partners)

DEVELOPER PERSONA

- Want customization and extensibility

- Expect APIs, SDKs, workflow builder tools

- Need documentation and developer support

- High technical fluency

- Measured by workflows created and API usage

Next-Gen Direction: Dual-Mode UX Strategy

DUAL-MODE UX STRATEGY

Mode 1: Curated Workflows (for Users)

Pre-built, one-click workflows designed by DeepJudge or trusted partners. Examples: "M&A Due Diligence," "NDA Review," "Regulatory Compliance Check." Users experience these as simple, guided processes with minimal configuration required.

Mode 2: Workflow Builder (for Builders)

Low-code/no-code workflow builder enabling legal engineers and partners (like Thomson Reuters) to create custom, multi-step AI applications. Builders can chain retrieval, analysis, generation, and external API calls into firm-specific or department-specific workflows.

Strategic Implication for the PM Role: The Lead Product Manager must balance both personas. Over-indexing on Users creates a feature-rich SaaS app but limits platform adoption. Over-indexing on Builders creates a powerful API but alienates non-technical end-users. Success requires a dual roadmap: curated workflows that demonstrate value to Users, and a workflow builder that empowers Builders to create the next generation of curated workflows. Thomson Reuters is the ultimate test case for the Builder persona.

Integration and Data Governance: System Integration

INTEGRATION ARCHITECTURE

- **DMS Connectors:** iManage, NetDocuments, Worldox, OpenText eDOCS

- **Collaboration Platforms:** SharePoint, Microsoft 365, HighQ

- **Email Archives:** Exchange, Outlook, Gmail

- **Custom APIs:** RESTful APIs for proprietary internal systems

- **Automated Sync:** Real-time or scheduled indexing with incremental updates

PRIVACY-BY-DESIGN PRINCIPLES

- **No Training on Client Data:** Client documents never used to train or fine-tune models

- **Data Residency Options:** EU, US, Switzerland, or on-premises deployment

- **Encryption:** AES-256 at rest, TLS 1.3 in transit

- **Audit Logging:** Complete audit trail of all queries, access, and system events

- **Data Minimization:** Only index what's necessary; configurable retention policies

Integration and Data Governance: Security Architecture

PERMISSION-AWARE INDEXING

- **Mirrors Source System ACLs:** Every document inherits original access control lists

- **Ethical Walls:** Automatically enforces Chinese walls and conflict-of-interest restrictions

- **Matter-Level Permissions:** Granular access control at matter, client, and document level

- **Dynamic Updates:** Permission changes in source systems propagate to DeepJudge in real-time

- **User-Centric Security:** Search results filtered based on authenticated user's permissions

SECURITY ARCHITECTURE

- **Flexible Deployment:** Cloud (AWS/Azure/GCP), private cloud, or on-premises

- **SOC 2 Type II:** Independently audited security controls

- **ISO 27001:** Information security management certification

- **GDPR Compliance:** Full compliance with EU data protection regulations

- **Penetration Testing:** Regular third-party security assessments

- **Zero-Trust Architecture:** Assume breach, verify every access request

---

The "Fort Knox" Positioning: This comprehensive security and governance architecture is not a feature checklist. It is a strategic weapon. It enables DeepJudge to win deals that cloud-only, permission-naive competitors cannot touch: Swiss banks, German industrial conglomerates, national security law firms, and any organization where data sovereignty and security are non-negotiable.

Role Definition: Platform Ecosystem PM

This is not a SaaS PM role adding features to a chat app, this is a PaaS PM role building a platform ecosystem

THE MANDATE

The Lead Product Manager **owns Layer 3** (AI Workflows platform). The mandate is to transform DeepJudge from a tool that lawyers use into a **platform that firms and partners build on**. This requires enabling "Builders" (legal engineers, partners like Thomson Reuters) to create custom, multi-step AI applications while simultaneously serving "Users" (lawyers) with curated, pre-built workflows. The role is responsible for fostering a developer ecosystem, managing platform APIs, and balancing the needs of two distinct user personas.

DUAL USERS

End-Users (Lawyers)

- Consume pre-built, curated workflows

- Expect simple, one-click UX

- Low technical fluency

- Measured by adoption and daily usage

Builders (Legal Engineers / Partners)

- Create custom workflows using platform tools

- Expect APIs, SDKs, documentation

- High technical fluency

- Measured by workflows created and API usage

Role Definition—Dual Stakeholders and Success Metrics

DUAL STAKEHOLDERS

DeepJudge (Direct Customers)

- Law firms and corporate legal departments
- Expect reliable, high-value workflows
- Revenue from SaaS subscriptions
- Success = adoption, retention, expansion

Thomson Reuters (Strategic Partner)

- Largest "Builder" and distribution partner
- Expect robust APIs and integration support
- Revenue from PaaS consumption and joint GTM
- Success = integration milestones, co-selling

SUCCESS METRICS

User Adoption

Daily active workflows, query volume, feature utilization across law firms

Builder Adoption

Workflows created by legal engineers, API call volume, developer satisfaction

TR Partnership

Integration milestones, joint GTM launches, co-selling pipeline

Platform Health

Ecosystem growth, third-party integrations, developer community engagement

Strategic Positioning: This role sits at the intersection of product, platform, and partnerships. It requires the technical depth to earn credibility with PhD founders and enterprise engineers, the strategic fluency to navigate a high-stakes partnership with Thomson Reuters, and the product instincts to balance two user personas with competing needs. This is not a role for a generalist PM, it requires platform expertise, enterprise fluency, and AI/ML technical depth.

Success Metrics and Strategic Imperatives

Concrete, measurable outcomes for the first 3-6 months

3-6 MONTH SUCCESS CRITERIA

PLATFORM LAUNCH
Ship V1 of AI Workflows platform to 3-5 pilot customers with core workflow builder functionality

BUILDER ADOPTION
Achieve 15+ custom workflows created by legal engineers across pilot firms

TR INTEGRATION
Complete Phase 1 technical integration: API connectivity, data unification, first joint workflow

USER ADOPTION
Maintain >70% daily active usage in pilot firms for curated workflows

PLATFORM STABILITY
Achieve <2% error rate, >99.5% uptime, <500ms p95 latency for workflow execution

ECOSYSTEM GROWTH
Onboard 5-10 "Builder" customers (legal engineers, partners) actively using workflow builder

STRATEGIC IMPERATIVES

Execute TR Partnership
Deliver on all integration commitments, establish joint governance model, align on roadmap priorities

Platform Foundation
Build workflow builder with intuitive low-code/no-code UX, robust API documentation, developer sandbox

Dual-Persona Balance
Ship curated workflows for Users while enabling Builders to create custom workflows, serve both simultaneously

Technical Excellence
Establish platform reliability, performance, and security standards that earn enterprise trust

Thomson Reuters Partnership KPIs and Measurement Framework

THOMSON REUTERS PARTNERSHIP KPIS

→ **Integration Milestones:** API integration complete, data unification operational, 3+ joint workflows deployed

→ **Joint GTM:** Co-selling pipeline established, joint marketing launches executed, partner enablement complete

→ **Technical Success:** >99.9% API uptime, <200ms API latency, <0.5% error rate for TR-facing endpoints

→ **Business Success:** Revenue from TR partnership on track, expansion opportunities identified and scoped

How Success Will Be Measured:

Monthly reviews with founders covering:
(1) Platform adoption metrics (Users and Builders),
(2) TR partnership health (integration milestones, joint GTM progress),
(3) Product quality (reliability, performance, developer satisfaction),
(4) Ecosystem growth (workflows created, API usage, developer community engagement). Success is not shipping features, it is building a platform that developers want to build on and users want to use.

The Dual Challenge: Dual Personas

Why this role requires exceptional strategic navigation and platform expertise

CHALLENGE 1: DUAL PERSONAS: THE BALANCING ACT

The Tension

Users (lawyers) want simplicity and pre-built workflows. Builders (legal engineers, TR) want power, flexibility, and APIs. These needs are **fundamentally in tension**.

Risk: Over-Index on Users

Creates a feature-rich SaaS app with great UX but **limits platform adoption**. Alienates Builders. TR partnership stalls. Platform potential unrealized.

Risk: Over-Index on Builders

Creates a powerful API and workflow builder but **poor end-user UX**. Low lawyer adoption. Direct customers churn. Revenue suffers.

The Required Balance: A dual roadmap that delivers curated workflows for Users (demonstrating immediate value and driving adoption) AND a workflow builder for Builders (enabling customization and ecosystem growth). Success requires serving both personas without compromising either.

The Dual Challenge: The TR Partnership

CHALLENGE 2: THE TR PARTNERSHIP, OPPORTUNITY AND RISK

The Opportunity

→ Distribution at scale through TR's global legal customer base

→ Validation from industry leader accelerates enterprise sales

→ Revenue acceleration via joint GTM and co-selling

→ Technical validation of platform capabilities

→ Ecosystem catalyst: TR success attracts other partners

The Risk

⚠ Dependency on single partner creates strategic vulnerability

⚠ TR roadmap influence could compromise product independence

⚠ Potential conflict: TR as both partner AND competitor

⚠ Resource allocation tension: TR demands vs. direct customers

⚠ If TR builds their own layer, DeepJudge becomes commoditized

The Required Balance: Maintain platform independence and product control while deepening the TR partnership. This requires clear contractual boundaries, a diversified customer base (so TR is not >50% of revenue), and a product roadmap that serves TR's needs without becoming TR-specific. The PM must be both partner advocate and platform guardian.

Why This Dual Challenge Matters: These are not implementation details, they are existential strategic questions. Get the persona balance wrong, and DeepJudge becomes either a niche developer tool or a feature-limited SaaS app. Get the TR partnership wrong, and DeepJudge either loses its biggest growth catalyst or becomes strategically dependent on a single partner.
This role requires a PM who can navigate these tensions with platform expertise, enterprise fluency, and strategic discipline.

AI Product Leadership: Vorwerk International

Senior Product Manager - Connected Products (IoT/AI)

Vorwerk International, Wollerau, Switzerland • 03/2016 - 04/2021 • CHF 100M Portfolio

PLATFORM VISION & PRODUCT STRATEGY

→ Owned global roadmap for connected and AI-enabled kitchen products integrated with Thermomix TM6, shaping strategy toward TM7

→ Created "Thermomix as a Service" platform combining connected hardware, personalized content, and subscription economics under privacy-by-design architecture

→ Filed patent for Multidevice Cooking, establishing Vorwerk's platform for synchronized smart-kitchen experiences and future AI orchestration

ML/AI IMPLEMENTATION & DATA PIPELINES

→ Defined telemetry and data pipelines that power ML-driven personalization and predictive maintenance across connected device ecosystem

→ Delivered GDPR-compliant AI solutions balancing personalization with privacy-by-design principles

→ Positioned AI as core value driver for next-generation product roadmap (TM7)

PRODUCT LAUNCHES & MARKET EXPANSION

→ Launched AI and IoT products (smart thermometer, smart pan) in UK, Spain, USA, and Australia

→ Opened new subscription revenue streams through connected product ecosystem

→ Orchestrated cross-functional teams across hardware, firmware, cloud, UX, Legal, Finance, and Compliance

CROSS-FUNCTIONAL LEADERSHIP

→ Led multicultural teams (120 engineers) across Switzerland, Germany, and Spain using agile methods (Scrum, SAFe)

→ Partnered with Legal, Compliance, and Finance to ensure GDPR compliance for all AI/IoT initiatives

→ Aligned hardware, firmware, cloud, and UX teams to deliver cohesive product experiences

MEASURABLE BUSINESS IMPACT

35%

Reduction in support incidents through predictive maintenance

20%

Lift in Salesforce customer engagement through connected products

CHF 100M

Portfolio value managed across global markets

AI Product Leadership: V-Zug AG

Digital Product Manager / AI Solutions Architect

V-Zug AG, Zug, Switzerland • 03/2022 - 05/2024 • 500k Connected Appliances

AI SOLUTIONS & PREDICTIVE ANALYTICS

→ Delivered predictive maintenance and subscription analytics for 500k connected appliances in production-scale deployment

→ Led AI initiatives from concept to production in regulated ecosystem, integrating business development and operational delivery

→ Defined, validated, and secured business cases for new digital and AI-powered services in hardware-centric environment

PRIVACY-BY-DESIGN & REGULATORY COMPLIANCE

→ Managed regulatory-compliant data flows and privacy-by-design integration for Swiss market, partnering with Legal, Compliance, and Finance for every new initiative

→ Ensured GDPR compliance and data residency controls for 500k connected devices across Switzerland and EU markets

→ Implemented privacy-by-design principles as foundational architecture for all AI and digital services

PRODUCT-AS-A-SERVICE INNOVATION

→ Drove product-as-a-service and digital business model innovation as digital product manager, acting as CEO of the product

→ Led agile transformation across cross-functional teams in hardware-centric organization

→ Integrated business development and operational delivery to drive digital transformation

ACADEMIC PURSUIT & APPLIED RESEARCH

→ Pursued MSc in Artificial Intelligence in parallel with full-time product leadership (2022-2024)

→ Applied latest research and trustworthy AI practices directly to commercial and technical challenges

→ Bridged academic rigor with production-scale AI deployment in regulated environment

SWISS MARKET EXPERTISE & REGULATORY MASTERY

500k

Connected appliances with privacy-by-design architecture

Production-Scale

AI deployment in regulated Swiss and EU markets

MSc AI

Completed in parallel with full-time product leadership

Platform Experience: L1/L2/L3 Stack Evolution

Building infrastructure, platform services, and applications at scale across Vorwerk and V-Zug

L1: Infrastructure Layer

VORWERK (2016-2021) + V-ZUG (2022-2024)

Built data pipelines, telemetry systems, and IoT connectivity infrastructure supporting millions of connected devices. Implemented privacy-by-design architecture as foundational layer for all AI and digital services. Established GDPR-compliant data flows and regulatory controls for Swiss and EU markets.

L2: Platform Services Layer

VORWERK (2016-2021) + V-ZUG (2022-2024)

Delivered ML-driven personalization and predictive maintenance services across connected device ecosystems. Built subscription analytics and predictive models for 500k+ appliances. Created platform services enabling product-as-a-service business models, reducing support incidents by 35%.

L3: Applications Layer

VORWERK (2016-2021) + V-ZUG (2022-2024)

Created "Thermomix as a Service" platform combining connected hardware, personalized content, and subscription economics. Drove product-as-a-service innovation in hardware-centric organizations. Filed patent for Multidevice Cooking, establishing platform for future AI orchestration.

INDEPENDENT AI RESEARCH & AGENTIC WORKFLOWS

11/2024 – PRESENT

→ Built tool-orchestration agent prototypes with experts using Python, Docker, and n8n to demonstrate end-to-end integration of assistant workflows, data pipelines, and inference services

→ Completed MSc AI thesis on LLM limitations and extended results into diagnostics for conversational model reliability and evaluation, supporting Responsible AI adoption

→ Engaged with ETH Zurich, LAC Lucerne, and open-source communities to share applied research on conversational AI and evaluation

Platform Experience: V-Zug and Cross-Functional Orchestration

PLATFORM EXPERIENCE AT SCALE

V-Zug
500K Connected Devices

→ Scaled platform to 500K connected appliances in Swiss and European markets

→ Enterprise B2B2C model: manufacturers → installers → homeowners

→ Multi-stakeholder orchestration with complex permission and data governance requirements

→ Privacy-by-design architecture for Swiss market compliance (GDPR, data residency)

CROSS-FUNCTIONAL ORCHESTRATION

Engineering
Earned credibility with ML/IoT engineers through technical depth and hands-on prototyping

Product
Balanced technical complexity with user needs, shipping consumer-grade UX on enterprise infrastructure

Enterprise
Navigated complex B2B2C relationships, multi-stakeholder alignment, and partnership execution

Compliance
Implemented GDPR compliance, data residency, and privacy-by-design from day one

Direct Relevance to DeepJudge: This experience maps directly to DeepJudge's platform challenge. The three-layer architecture (Infrastructure → Platform → Applications) mirrors DeepJudge's evolution (Search → Assistant → Workflows). The dual-persona experience (end-users and builders) is exactly the challenge of serving lawyers and legal engineers. The B2B2C orchestration at V-Zug parallels the DeepJudge-Thomson Reuters partnership complexity. The privacy-by-design expertise is essential for DeepJudge's "Fort Knox" positioning.

Technical Depth: Academic Foundation

MSc in Artificial Intelligence with research on AGI architectures and LLM limitations

MSc in Artificial Intelligence

Institution: Aalen University, Germany

Dates: 2021–2024

Grade: 1.3/1.0 (top honors)

Thesis: "On Artificial General Intelligence" (graded 1.0/1.0)

Parallel Achievement: Completed MSc in parallel with full-time product leadership at V-Zug (2022-2024), applying latest research directly to commercial AI challenges

Thesis: "On Artificial General Intelligence"

→ Analyzed transformer limits, world-model architectures, interpretability, and safe alignment strategies

→ Investigated AGI architectures and world models, integrating LeCun's AMI concepts to explore world-inference beyond transformers

→ Developed AGI risk framework covering alignment, controllability, and ethical safeguards

→ Emphasized risk management crucial for conversational AI and production deployment

RESEARCH: "OPTIMIZING LLM ARCHITECTURES, ADDRESSING PROMPT ENGINEERING LIMITATIONS"

→ Critiqued prompt engineering's reliability by exposing structural divergences in tokenization, architecture, and scaling, showing prompt variability as an inherent model limitation

→ Proposed design changes that integrate user-interaction constraints into model and tokenizer design so promptability is built in rather than post hoc

→ Extended thesis results into diagnostics for conversational model reliability and evaluation, supporting Responsible AI adoption

RELEVANCE TO DEEPJUDGE'S LEGAL AI PLATFORM

Understanding **LLM limitations and prompt engineering reliability** is critical for building trustworthy legal AI. My research on transformer limits, world-model architectures, and prompt variability directly addresses the challenges DeepJudge faces in building reliable, production-grade legal workflows. The AGI risk framework and Responsible AI focus align with DeepJudge's need for **Fort Knox architecture and enterprise trust**. This academic depth enables technical conversations with PhD founders (ETH Zurich, Google Research) as peers, not just product execution.

Technical Depth: Production ML Systems

Hands-on ML engineering experience across computer vision, NLP, and industrial AI

COMPUTER VISION: IDENTITY VERIFICATION

→ Designed CNN-pipelines using ResNet50V2, MobileNetV2, and NASNetLarge for scalable identity validation

→ Applied transfer learning and callbacks to optimize accuracy under resource constraints

→ Production-ready architecture for regulated identity verification use cases

NLP: LARGE-SCALE TEXT ANALYSIS

→ Processed 7.5 GB corpus for token-frequency analytics and semantic pattern extraction

→ Prototyped Python pipelines for efficient document processing and transaction-record analysis

→ Demonstrated scalability for enterprise-grade text analytics

INDUSTRIAL AI: MANUFACTURING INTELLIGENCE

→ Built ML pipelines to detect rare pattern anomalies, improving quality control accuracy by 20%

→ Deployed predictive maintenance models using IoT data, reducing service intervals by 15%

→ Production deployment in regulated manufacturing environment

INDEPENDENT LLM & AGENTIC WORK (2024-PRESENT)

→ Built tool-orchestration agent prototypes using Python, Docker, and n8n platform

→ Completed "Building AI Agents" course (dair.ai), deploying fully autonomous AI agents

→ Prompt engineering, LLM prototyping, and conversational AI evaluation

TECHNICAL STACK & TOOLS

ML Frameworks
Python, TensorFlow, PyTorch

Infrastructure
Docker, Git, containerized ML workflows

AI & Data
Prompt engineering, LLM prototyping, ML feature design, statistical modeling

Analytics & Visualization
Analytics dashboards, data visualization, model evaluation

Cloud & Platforms
GCP, Azure, AWS, API integration

Product & Delivery
Agile (Scrum, SAFe), Jira, Confluence, strategic roadmapping

Regulatory Mastery: GDPR and Enterprise Compliance

Privacy-by-design expertise essential for DeepJudge's Fort Knox security architecture

PRIVACY-BY-DESIGN

✓ Implemented at **V-Zug** (500K connected devices)

✓ Privacy-first architecture from day one

✓ Data minimization and purpose limitation

✓ User consent and transparency

✓ Privacy impact assessments (PIAs)

GDPR & DATA GOVERNANCE

✓ **GDPR compliance** for Swiss and EU markets

✓ Data residency controls (Swiss data in Switzerland)

✓ Right to access, rectification, erasure (Articles 15-17)

✓ Data processing agreements (DPAs) with third parties

✓ Audit trails and compliance documentation

Regulatory Mastery: DeepJudge Fort Knox Relevance

DEEPJUDGE FORT KNOX ARCHITECTURE, DIRECT RELEVANCE

Permission-Aware Architecture

DeepJudge's core differentiator is zero data leakage through permission-aware indexing. I implemented similar multi-tenant, permission-based systems at V-Zug (B2B2C model with complex access controls).

Swiss Data Residency

DeepJudge positions as Swiss-first with data residency controls. I have deep Swiss market experience (V-Zug) and understands regulatory requirements for Swiss enterprise customers.

Enterprise Trust

Law firms demand the highest security standards, DeepJudge's Fort Knox positioning is a competitive advantage. I have enterprise compliance fluency to maintain and communicate this trust.

Regulatory Compliance

Legal industry is highly regulated. DeepJudge must navigate GDPR, data protection laws, and professional secrecy rules. My GDPR expertise is essential for product decisions.

Why This Expertise Is Critical: DeepJudge's Fort Knox architecture is not just a technical feature, it is the foundation of enterprise trust and a core competitive advantage against ChatGPT and Perplexity. Building Layer 3 (AI Workflows) requires a PM who deeply understands privacy-by-design, can navigate GDPR compliance, and can communicate security guarantees to enterprise customers. This is not checkbox compliance, it is strategic differentiation. My privacy-first experience at V-Zug, combined with Swiss market fluency, makes me uniquely qualified to protect and extend DeepJudge's Fort Knox positioning.

90-Day Action Plan: Discovery, Prioritization, Delivery

A concrete execution roadmap from day one to shipping V1 workflows platform

Phase 1: Discovery & Alignment

DAYS 1-30

Internal alignment, stakeholder mapping, technical deep dive, customer discovery

KEY ACTIVITIES

- Meet with founders: vision, priorities, constraints
- Technical deep dive: L1/L2/L3 architecture
- Interview 5-10 pilot customers (lawyers + legal engineers)
- TR partnership: integration commitments, timeline
- Team alignment: engineering, design, sales

DELIVERABLES

→ Stakeholder map with decision authority
→ Technical architecture documentation
→ Customer insights report
→ TR partnership roadmap alignment

Phase 2: Roadmap Prioritization

DAYS 31-60

Define V1 scope, prioritize features, establish metrics, create PRDs

KEY ACTIVITIES

- Define V1 scope: curated workflows + workflow builder
- Prioritize features based on discovery and TR needs
- Establish success metrics (adoption, builder usage)
- Create PRDs for top workflows and builder MVP
- Align with engineering on feasibility and timeline

DELIVERABLES

→ V1 product roadmap with prioritized features
→ PRDs for top 3 workflows and builder MVP
→ Success metrics dashboard framework
→ Engineering capacity plan and sprint roadmap

Phase 3: Execution & Shipping

DAYS 61-90

Ship V1 to pilot customers, onboard Builders, iterate based on feedback

KEY ACTIVITIES

- Ship V1 workflows platform to 3-5 pilot customers
- Onboard first Builder customers (legal engineers)
- Collect feedback and iterate on UX/functionality
- Begin TR integration Phase 1 (API connectivity)
- Establish weekly metrics review with founders

DELIVERABLES

→ V1 platform live with 3-5 pilot customers
→ 5+ workflows created by Builders
→ User feedback report with iteration priorities
→ TR integration Phase 1 milestone complete
→ 90-day retrospective and Q2 roadmap

90-DAY SUCCESS CRITERIA

✓ V1 platform shipped to pilot customers

✓ TR partnership on track with Phase 1 complete

✓ Team aligned and executing with velocity

✓ Positive feedback from Users and Builders

✓ Clear roadmap for next 6 months

✓ Founders confident in platform direction

12-Month Roadmap: Workflows to Agentic AI

Strategic progression from TR partnership launch to platform leadership and market maturity

Q1: TR Partnership Launch

MONTHS 1-3

Ship V1 workflows platform, execute TR integration Phase 1, establish platform foundation

KEY MILESTONES

→V1 workflows platform live with 3-5 pilot customers

→TR integration Phase 1 complete (API connectivity, data unification)

→15+ workflows created by Builders (legal engineers)

→Workflow builder MVP with low-code/no-code UX

SUCCESS METRICS

✔ >70% daily active usage, <2% error rate, 5-10 Builder customers onboarded

Q2: Platform Scaling

MONTHS 4-6

Scale to 10+ customers, deepen TR partnership, expand workflow library and ecosystem

KEY MILESTONES

→Scale to 10+ law firms (mix of direct and TR-sourced)

→TR integration Phase 2 (joint workflows, co-selling launch)

→Expand workflow library to 20+ curated workflows

→Developer documentation and API sandbox live

→First third-party integration (beyond TR)

SUCCESS METRICS

✔ 50+ workflows created by ecosystem, >80% adoption rate in customer base

12-Month Roadmap: Q3, Q4, and Strategic Vision

Q3: Agentic AI & Advanced Workflows

MONTHS 7-9

Ship agentic workflows, multi-step automation, expand to corporate legal departments

KEY MILESTONES

→Agentic workflows: autonomous multi-step legal tasks

→Advanced AI: reasoning, planning, tool use (beyond RAG)

→Corporate legal departments expansion (new segment)

→Platform API v2 with enhanced capabilities

→5+ third-party integrations live

SUCCESS METRICS

✓ 100+ workflows in ecosystem, agentic workflows in production use

Q4: Market Leadership

MONTHS 10-12

Establish market leadership, platform maturity, prepare for exponential scale

KEY MILESTONES

→25+ enterprise customers (law firms + corporate legal)

→TR partnership: joint GTM at scale, strong co-selling pipeline

→Platform ecosystem: 10+ third-party integrations

→Agentic AI workflows: 20+ autonomous workflows live

→Market positioning: recognized platform leader in legal AI

SUCCESS METRICS

✓ Platform revenue growing, ecosystem thriving, market leadership established

12-MONTH VISION: STRATEGIC TRANSFORMATION

→**From tool to platform:** DeepJudge becomes infrastructure layer for legal AI

→**From workflows to agents:** Autonomous AI handling complex legal tasks

→**From search to platform:** Complete L1 → L2 → L3 evolution realized

→**From TR to ecosystem:** Multiple partners building on DeepJudge

→**From Swiss to global:** Expansion beyond Switzerland to EU and beyond

→**Market positioning:** Platform leader, not feature competitor to ChatGPT

GTM Strategy: TR Partnership and Corporate Expansion

A dual-track approach balancing partnership leverage with strategic independence

TRACK 1: DEEPEN TR PARTNERSHIP

Opportunity: Thomson Reuters' global legal customer base provides distribution at scale and enterprise validation

GTM APPROACH

→ Joint value proposition: DeepJudge workflows integrated into TR ecosystem

→ Co-selling motion: TR sales team selling DeepJudge workflows

→ Joint marketing: Co-branded content, webinars, case studies

→ Technical integration: Seamless UX within TR products

→ Partner enablement: Train TR sales/support teams

SUCCESS METRICS

TR-sourced revenue, joint customer wins, integration milestones, co-selling pipeline growth

RISK MITIGATION

Maintain platform independence, ensure TR is not >50% of revenue, diversify customer base

GTM Strategy: Corporate Expansion and Strategic Rationale

TRACK 2: CORPORATE LEGAL EXPANSION

Opportunity: Corporate legal departments (in-house counsel) are underserved and provide revenue diversification

GTM APPROACH

→ Direct sales: Build enterprise sales team for corporate segment

→ Use cases: Contract review, compliance workflows, legal research

→ Positioning: Platform for legal operations teams (not just lawyers)

→ Pilot programs: 3-5 corporate legal departments in Q2-Q3

→ Case studies: Demonstrate ROI and efficiency gains

SUCCESS METRICS

Corporate customer wins, ARR from corporate segment, pilot-to-paid conversion rate

STRATEGIC VALUE

Reduces dependency on TR, diversifies revenue, validates platform with direct customers

STRATEGIC RATIONALE, WHY THIS DUAL APPROACH

- TR partnership provides distribution at scale and enterprise validation
- Both tracks serve the dual-persona model (Users + Builders)
- Reduces single-partner dependency risk while maximizing TR opportunity

- Corporate expansion provides revenue diversification and strategic independence
- Platform approach enables serving both segments with same infrastructure
- Corporate segment validates product-market fit beyond TR ecosystem

12-MONTH GTM MILESTONES

Q1 2025	Q2 2025	Q3 2025	Q4 2025
TR integration Phase 1 complete, joint GTM launch	First 3 corporate legal department pilots launched	TR co-selling motion operational, 5+ joint customer wins	Corporate segment contributing 20-30% of new ARR

Business Model Innovation: Dual Monetization

Platform economics enabling exponential scale through SaaS and PaaS revenue streams

SAAS MODEL FOR USERS

TARGET: LAWYERS AT LAW FIRMS AND CORPORATE LEGAL DEPARTMENTS

PRICING MODEL

→Starter: Basic workflows, limited queries ($99/user/month)

→Professional: Full workflow library, unlimited queries ($299/user/month)

→Enterprise: Custom workflows, priority support, SLA (custom pricing)

VALUE PROPOSITION

Pre-built curated workflows, simple one-click UX, immediate value, no technical expertise required

REVENUE DRIVERS

Seat-based pricing, expansion within firms, enterprise contracts, net revenue retention

SUCCESS METRICS

ARR growth, net revenue retention >110%, customer lifetime value, seats per customer

PAAS MODEL FOR BUILDERS

TARGET: THOMSON REUTERS, LEGAL TECH PARTNERS, LEGAL ENGINEERS

PRICING MODEL

→API calls: $0.01-0.05 per call (volume discounts)

→Workflow executions: $1-5 per execution

→Data processing: $0.10 per GB processed

→Enterprise contracts: Custom pricing with committed spend

VALUE PROPOSITION

Powerful APIs, workflow builder, developer tools, platform flexibility, white-label capabilities

REVENUE DRIVERS

API usage growth, workflow creation volume, platform adoption, TR partnership leverage

SUCCESS METRICS

API call volume, revenue per Builder, platform usage, workflows created, ecosystem growth

Business Model Innovation:
Strategic Rationale and Platform Economics

STRATEGIC RATIONALE, WHY DUAL MONETIZATION ENABLES PLATFORM SCALE

- SaaS provides predictable recurring revenue and direct customer relationships

- Network effects: More Builders → more workflows → more User value → more Builders

- Margin expansion: PaaS has higher gross margins (85%) than SaaS (70%)

- PaaS provides exponential scale through ecosystem leverage and network effects

- TR partnership: PaaS model aligns with TR's distribution and white-label needs

- Competitive moat: Platform with dual monetization is harder to replicate than SaaS-only

- Platform economics: 1 Builder can serve 100+ Users with custom workflows

- Revenue diversification: Not dependent on single revenue stream or customer segment

- Exponential scale: PaaS revenue grows faster than SaaS as ecosystem matures

PLATFORM ECONOMICS, REVENUE EVOLUTION AND MARGIN PROFILE

Year 1 Mix
70% SaaS, 30% PaaS (building foundation and TR integration)

Year 2 Mix
60% SaaS, 40% PaaS (ecosystem growth and Builder adoption)

Year 3 Mix
50% SaaS, 50% PaaS (platform maturity and exponential scale)

Platform Leverage
1 Builder creates workflows serving 100+ Users (10-100x leverage)

Margin Profile
SaaS 70% gross margin, PaaS 85% gross margin (lower support costs)

Growth Rate
PaaS revenue grows 2-3x faster than SaaS as ecosystem scales

Closing Argument & Next Steps

You have built the industry's most powerful **"Knowledge Layer."**
The market has validated your infrastructure-first thesis through elite firms like Gunderson Dettmer and the strategic Thomson Reuters partnership.

The **"Next-Gen AI" PM role** is the critical link to your platform future. This "0-to-1" opportunity requires a leader fluent in two specific domains:

THE CRITICAL NEED

→ **Full-Stack Platform Management:** The ability to manage a "dual-user" (L1-L3) product ecosystem

→ **"Fort Knox" Security & Compliance:** The credibility to build for a regulated, high-stakes enterprise and Swiss-based market

THE DIRECT MATCH

✓ **Platform Management:** I owned the global roadmap for a CHF 100M+ connected platform at Vorwerk, creating the "Thermomix as a Service" ecosystem

✓ **"Fort Knox" Compliance:** I led the delivery of GDPR-compliant AI solutions and managed regulatory-compliant data flows for 500k+ connected appliances at V-Zug

I have the plan, the technical depth, and the specific enterprise-scale platform experience to execute this role from day one.

Thank you for your time. I look forward to our discussion.

Marcelo Caballero • AI Product Manager • msc_ai@icloud.com